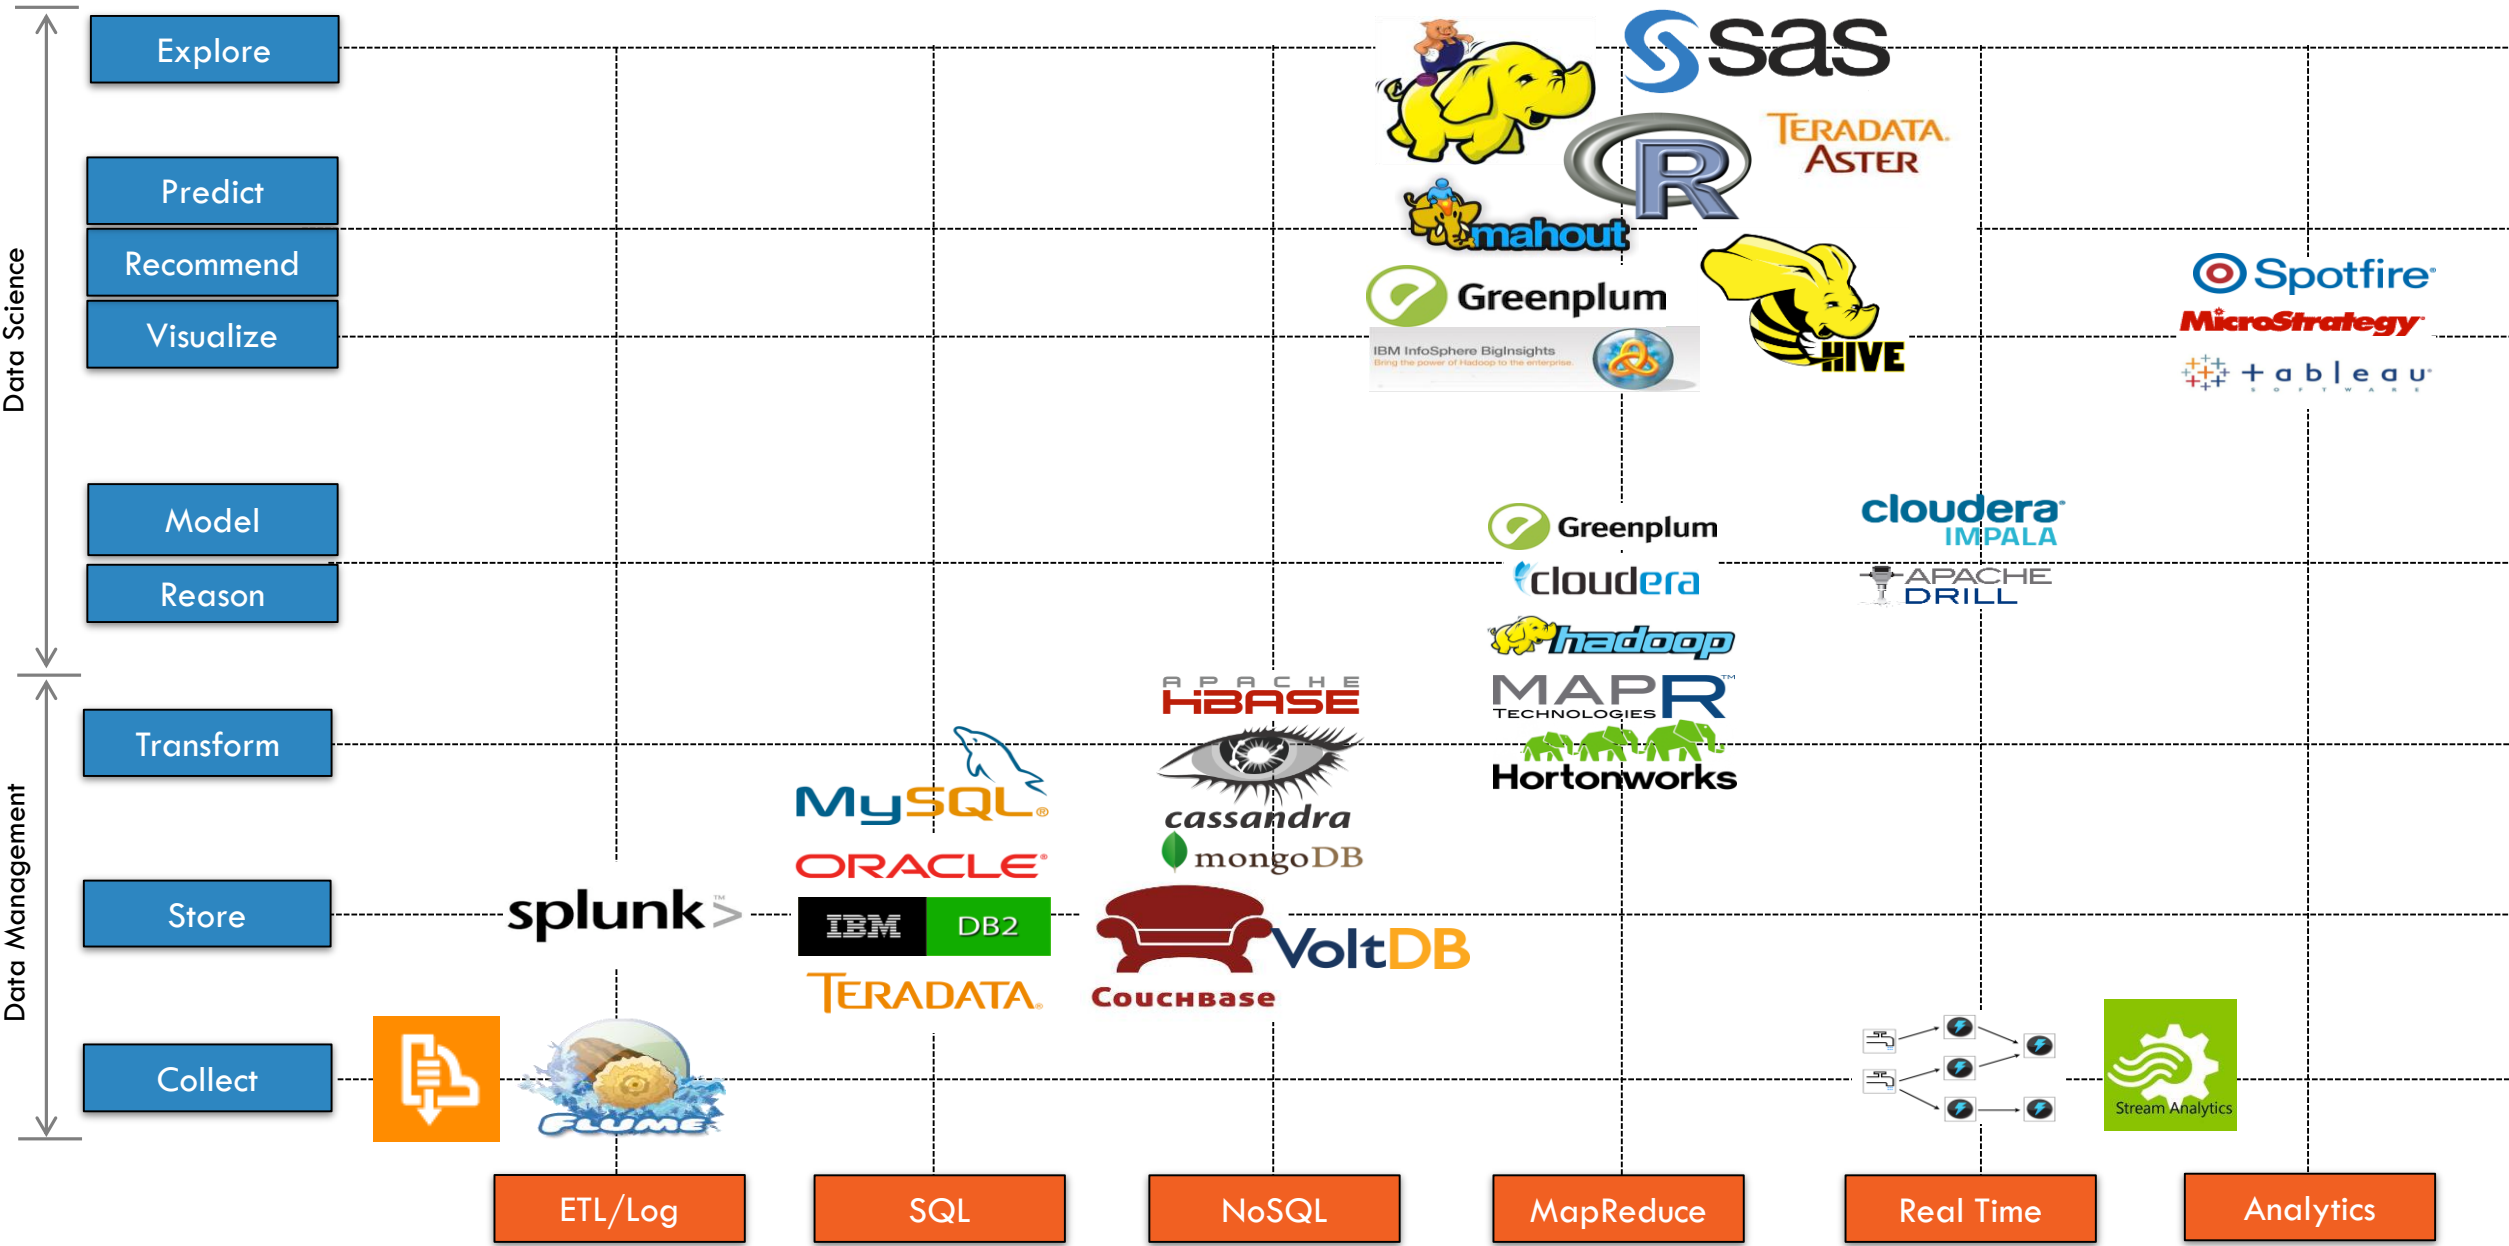


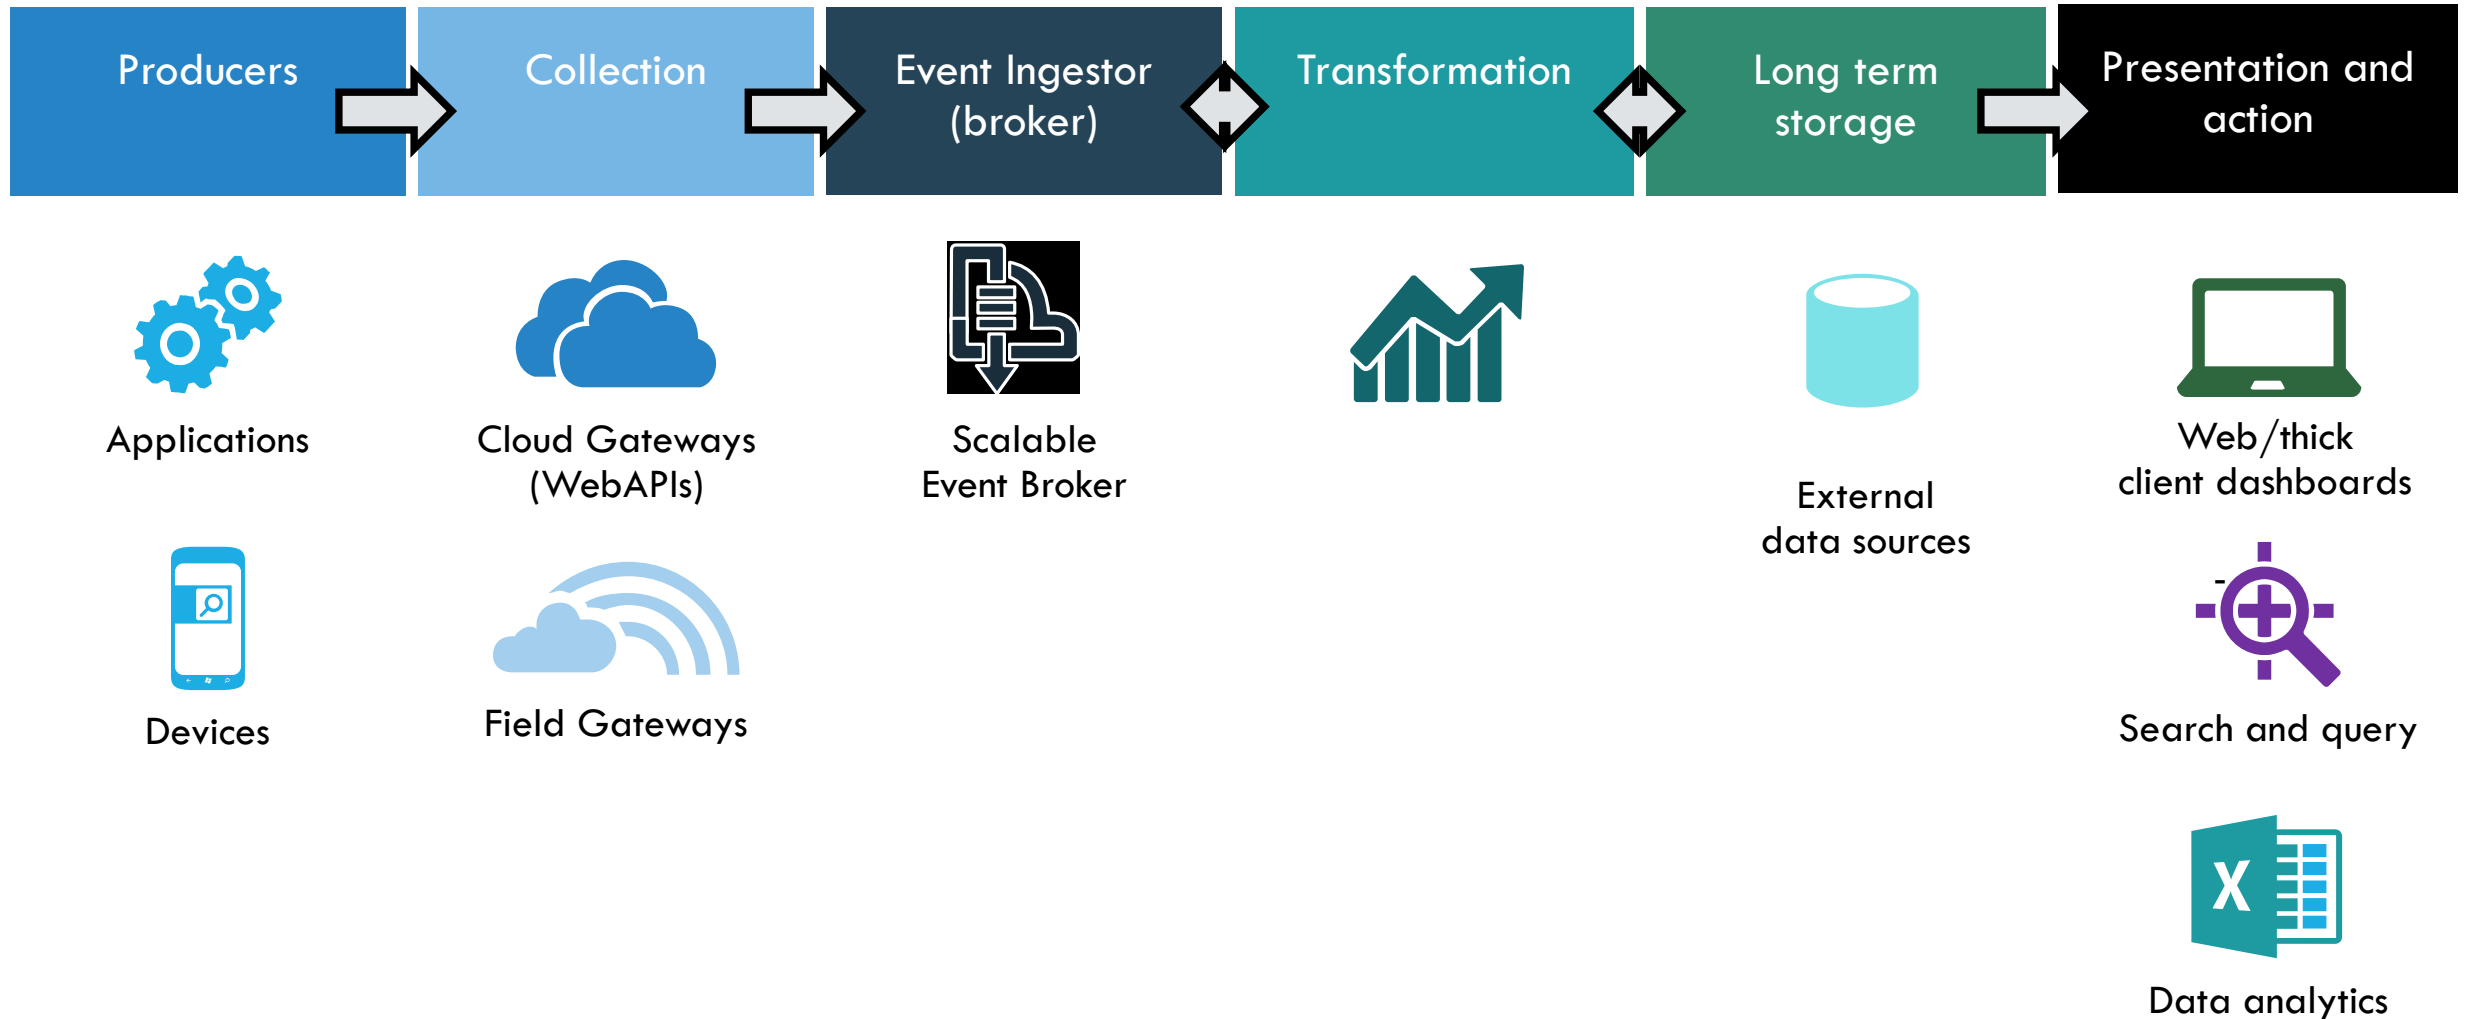


# EVENT PROCESSING WITH STREAM ANALYTICS



Big Data – Technology, Platforms & Products

# TYPICAL EVENT PROCESSING







1  
**NEW**  
DEFINITION  
IS ADDED ON  
urban

1,600+  
**READS ON**  
Scribd

13,000+ HOURS  
**MUSIC**  
STREAMING ON  
PANDORA

12,000+  
**NEW ADS**  
POSTED ON  
craigslist

370,000+ MINUTES  
**VOICE CALLS ON**  
skype

98,000+  
**TWEETS**

320+  
**NEW**  
twitter  
ACCOUNTS

100+  
**NEW**  
Linked in  
ACCOUNTS

1 associated content  
**NEW**  
ARTICLE IS  
PUBLISHED

THE  
WORLD'S  
LARGEST  
COMMUNITY  
CREATED CONTENT!

6,600+  
**NEW**  
PICTURES ARE  
UPLOADED ON  
flickr

50+  
**WORDPRESS**  
DOWNLOADS

695,000+  
**facebook**  
STATUS  
UPDATES

125+  
**PLUGIN**  
DOWNLOADS

79,364  
**WALL**  
POSTS

510,040  
**COMMENTS**

1,700+  
**Firefox**  
DOWNLOADS

694,445  
**SEARCH**  
QUERIES

168 MILLION  
**EMAILS**  
ARE SENT

60+  
**NEW**  
BLOGS

1,500+  
**BLOG**  
POSTS

70+  
**DOMAINS**  
REGISTERED

600+  
**NEW**  
VIDEOS

QUESTIONS  
ASKED ON THE  
INTERNET...

25+ HOURS  
**TOTAL**  
DURATION

# TIMELINESS OF INFORMATION

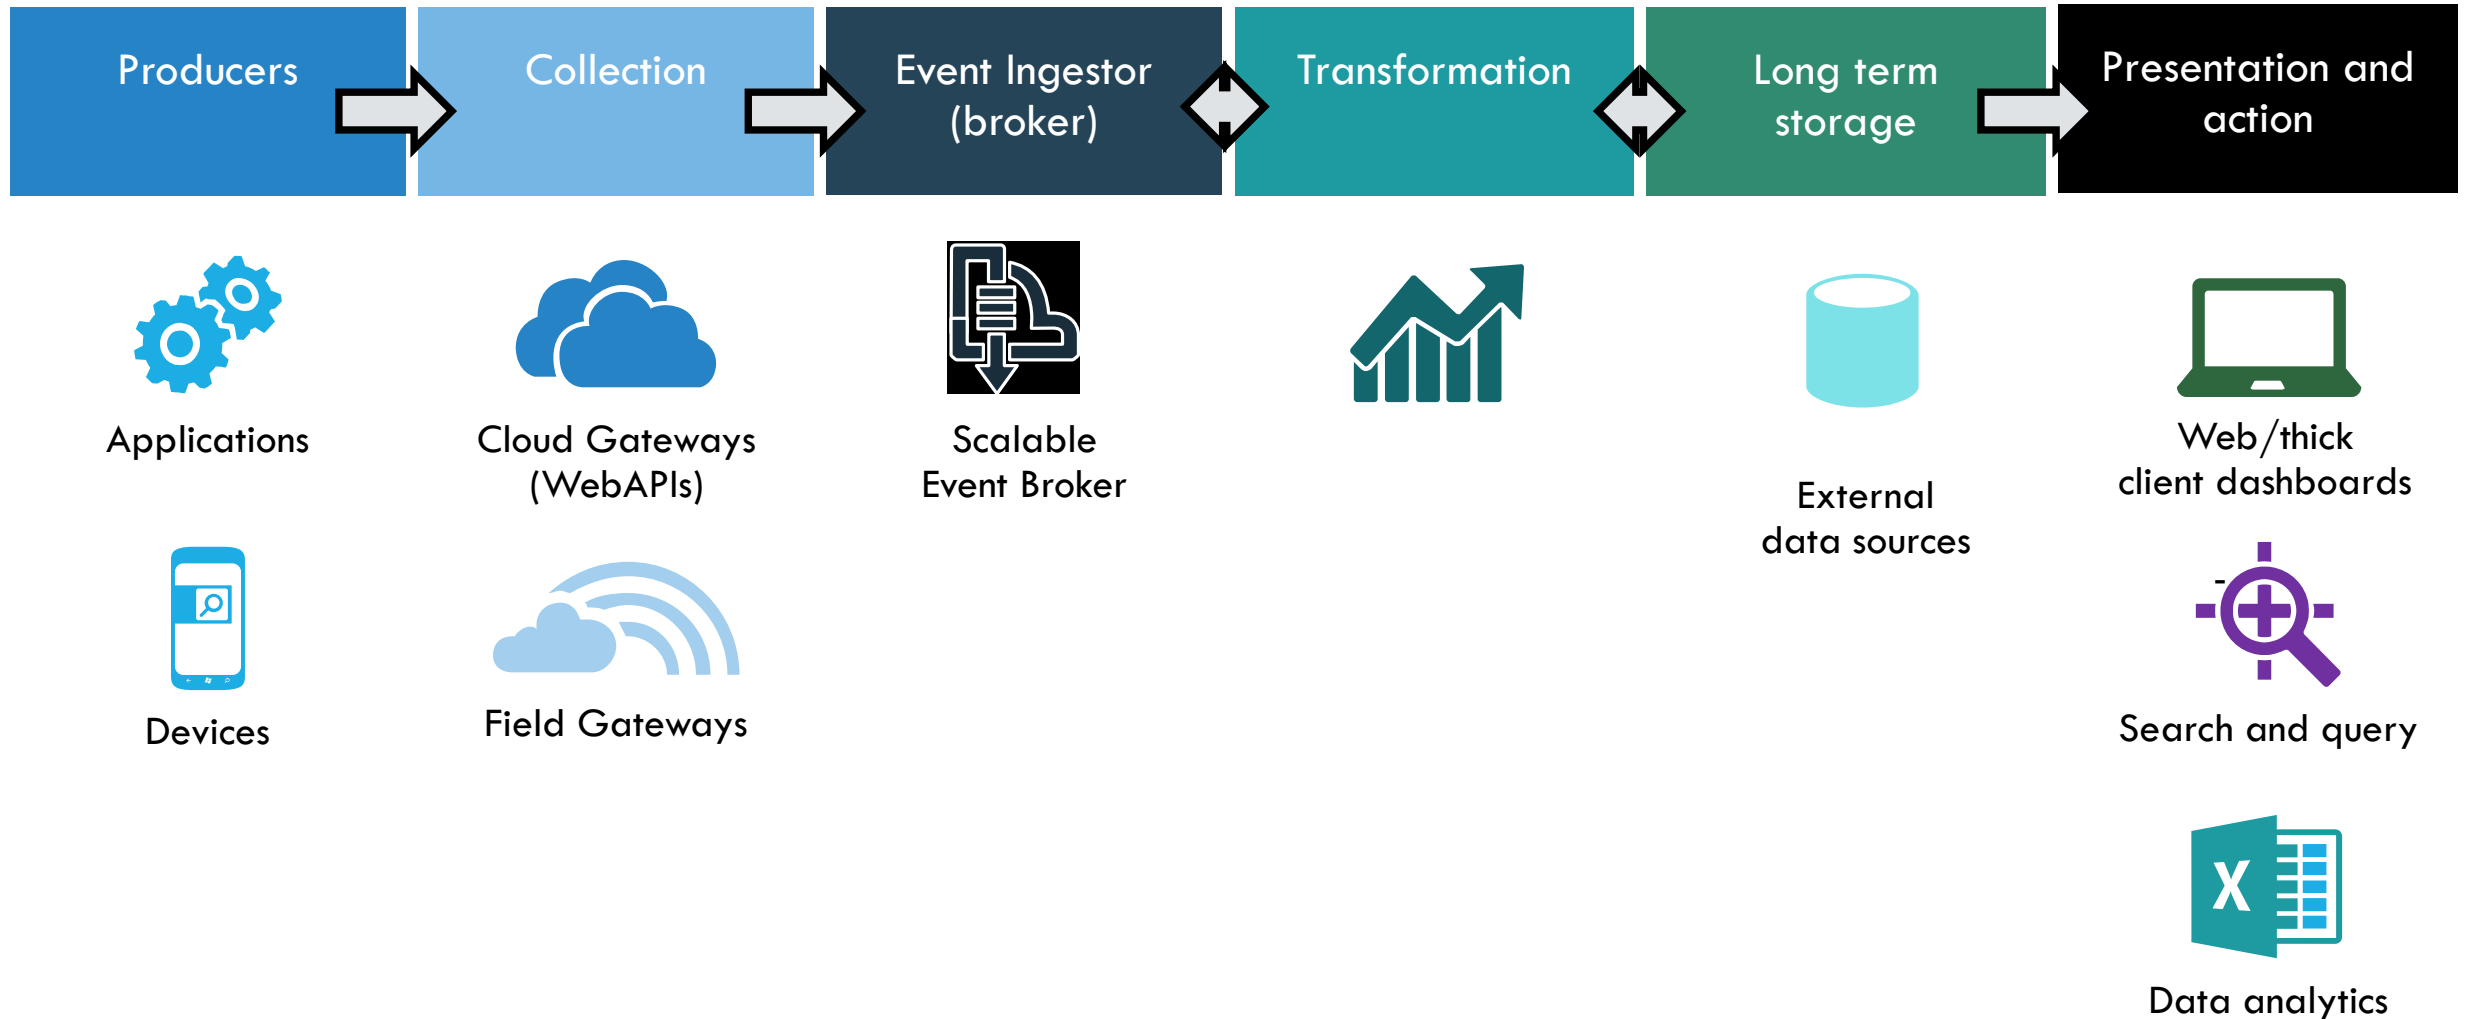


What's trending in the past 5 minutes?

Your high school friend is also in Vegas  
**NOW.**



# TYPICAL EVENT PROCESSING





# DATA AT REST

## Question

"How many red cars are in the parking lot?"

## Answering with a relational database

Walk out to the parking lot

Count vehicles that are: Red, Car

```
SELECT count(*) FROM ParkingLot  
WHERE type = 'Auto'  
      AND color = 'Red'
```



# DATA IN MOTION

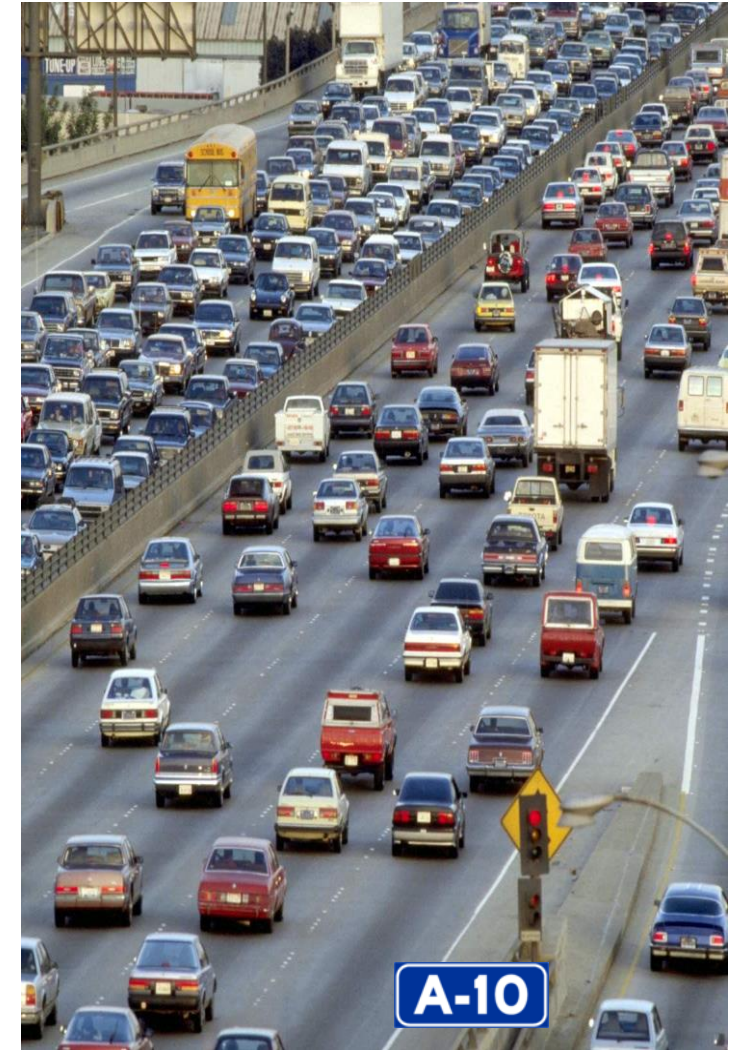
## Different Question

"How many red cars have passed exit 18A on A-10 in the last hour?"

## Answering with a relational database

Pull over, park all vehicles in a lot, keep them there for an hour  
Count vehicles in the lot

Not a great solution...





# AZURE STREAM QUERY LANGUAGE

Simple SQL dialect

**Familiar** – learning curve reduction

**High-Level** – expression of intent, not implementation

**Maintainable** – focus on the essentials of the problem

Extended in natural ways  
to express temporal concepts

WINDOW – multiple kinds

(tumbling, hopping, sliding)

TIMESTAMP BY, BETWEEN

DATEDIFF in joins

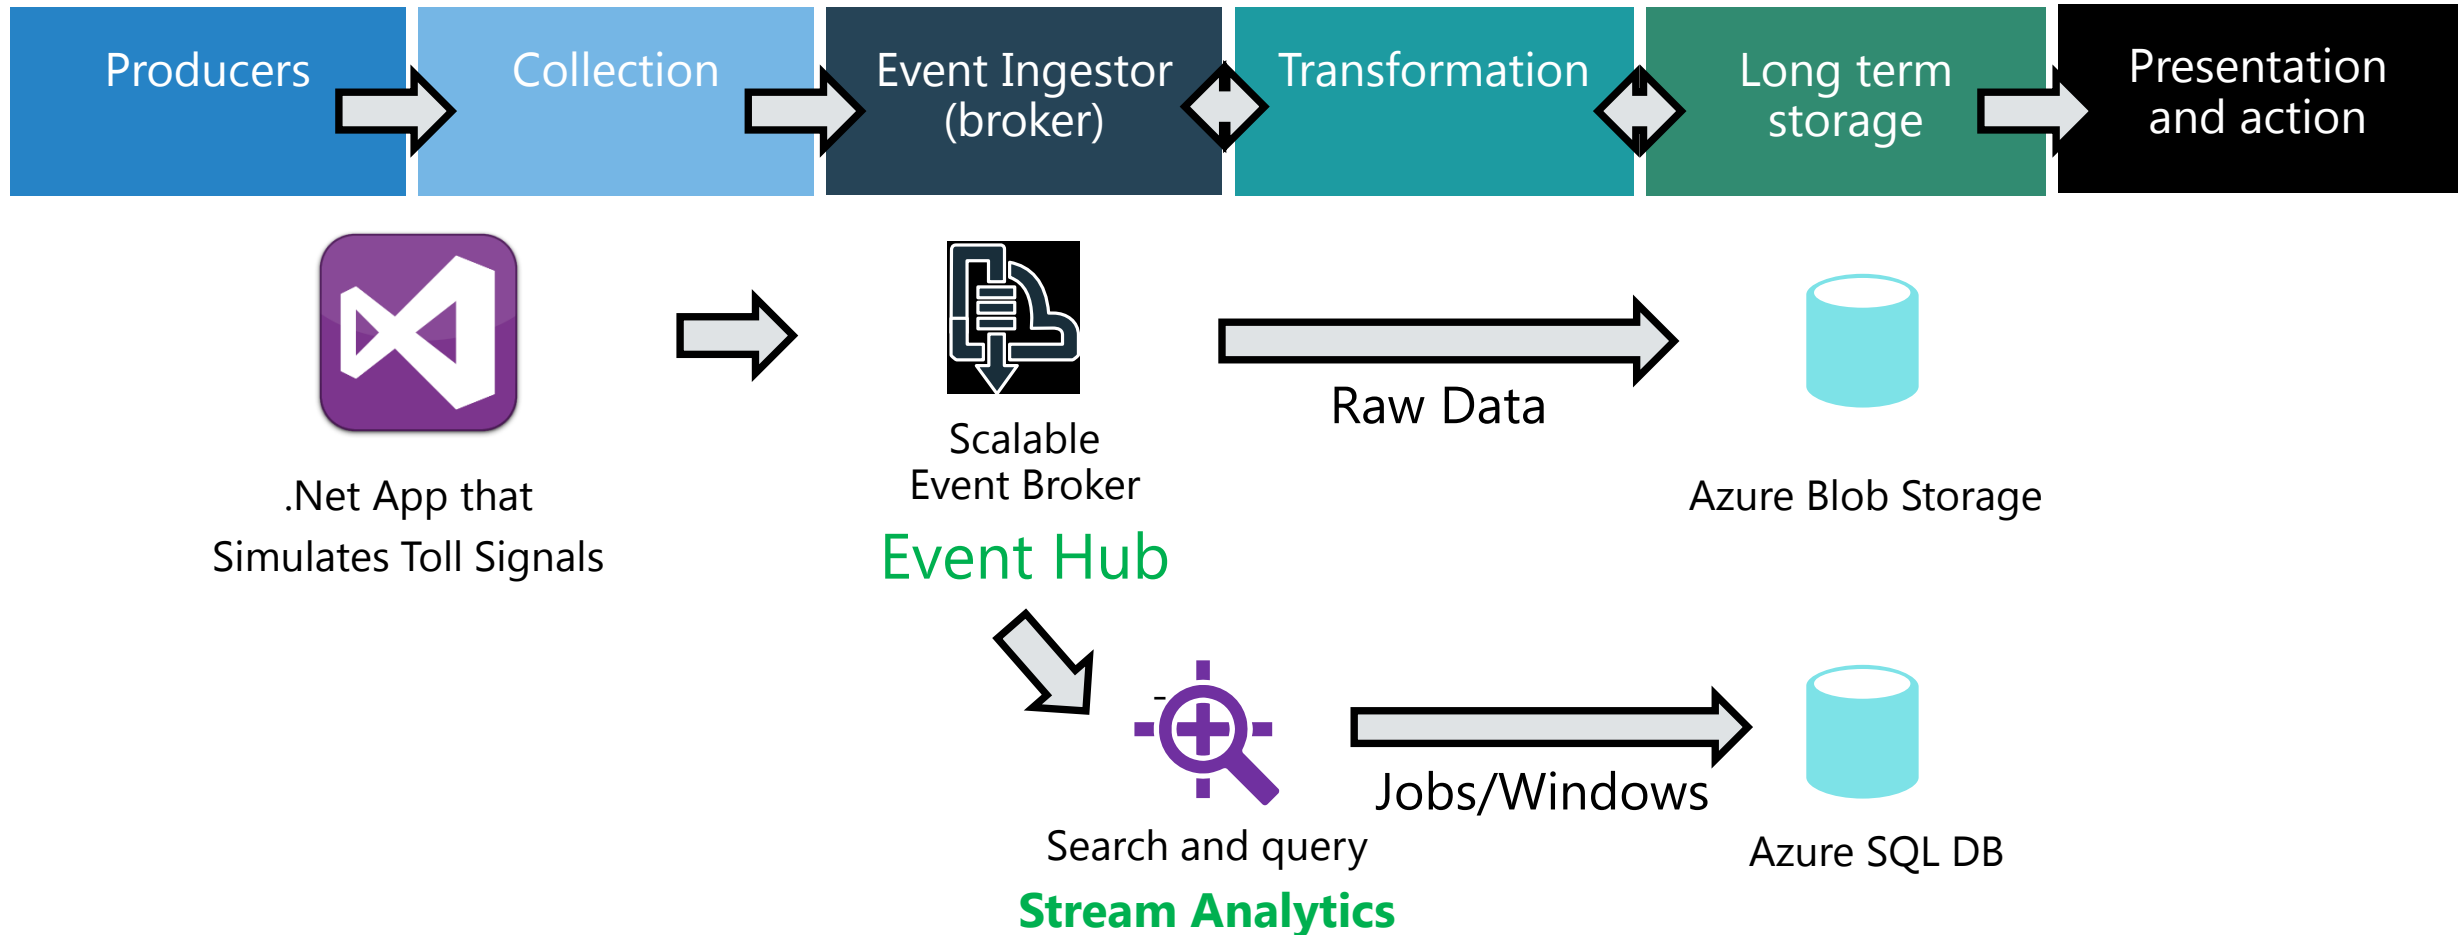
PARTITION BY for scale-out

```
WITH agg AS
(
    SELECT Avg(reading), Building
    FROM Temperature
    GROUP BY TumblingWindow(second, 1), building
)
SELECT A1.Avg AS Old, A2.Avg AS New, A1.Building
FROM Agg A1 JOIN Agg A2
ON A1.Building = A2.Building
AND DATEDIFF(minute,A1,A2) BETWEEN 4.5 AND 5.5
WHERE
    (a1.avg < a2.avg - 10) OR (a1.avg > a2.avg+10)
```



# DEMO

# TYPICAL EVENT PROCESSING





# STRENGTHS

- Analyze millions of events per SECOND
- Fault tolerant
- SQL spoken here
- Fully managed service by Azure

# BUILT-IN FUNCTIONS AND SUPPORTED TYPES

## Aggregate functions

Count, Min, Max, Avg, Sum

## Scalar functions

Cast

**Date and time:** Datename, Datepart, Day, Month, Year, Datediff, Dateadd

**String:** Len, Concat, Charindex, Substring, Patindex

# A TEMPORAL SYSTEM

- Every event is a point in time, and thus must come with a timestamp.
  - (remember how relational DBs need a PK? Temporal systems need a timestamp)
- Stream Analytics can append your events with a timestamp. (bad practice if standalone)
  - Can be skewed by network and hardware latency.
- Users can define application time stamps with the **TIMESTAMP BY** clause.
- Aggregations have timestamps at the end of the window.



# TRADITIONAL SQL

How many vehicles passed through each toll booth yesterday?

- Why can't we ask how many cars have gone through so far today?

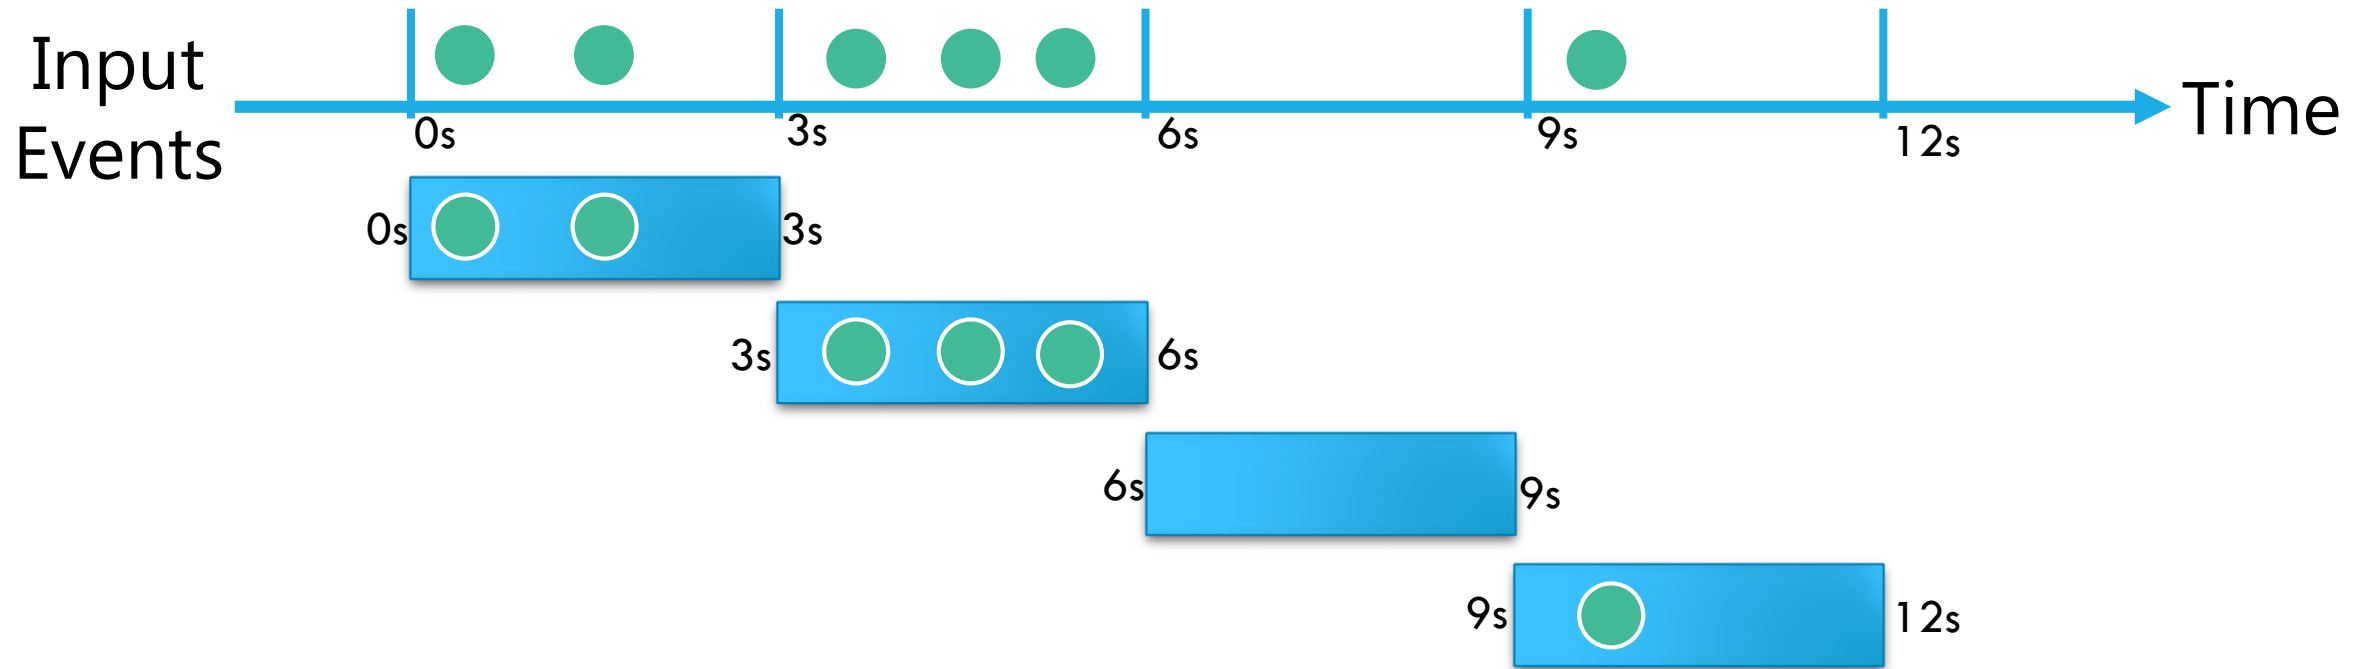
```
SELECT TollID, Count(*) AS Count
FROM EntryTable
WHERE date = 'yesterday'
GROUP BY TollID
```

# AZURE STREAM QUERY LANGUAGE

How many vehicles pass through each toll booth every 3 seconds?

```
SELECT TollID, System.Timestamp AS WindowEnd, Count(*) AS Count  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TUMBLINGWINDOW(second, 3), TollID
```

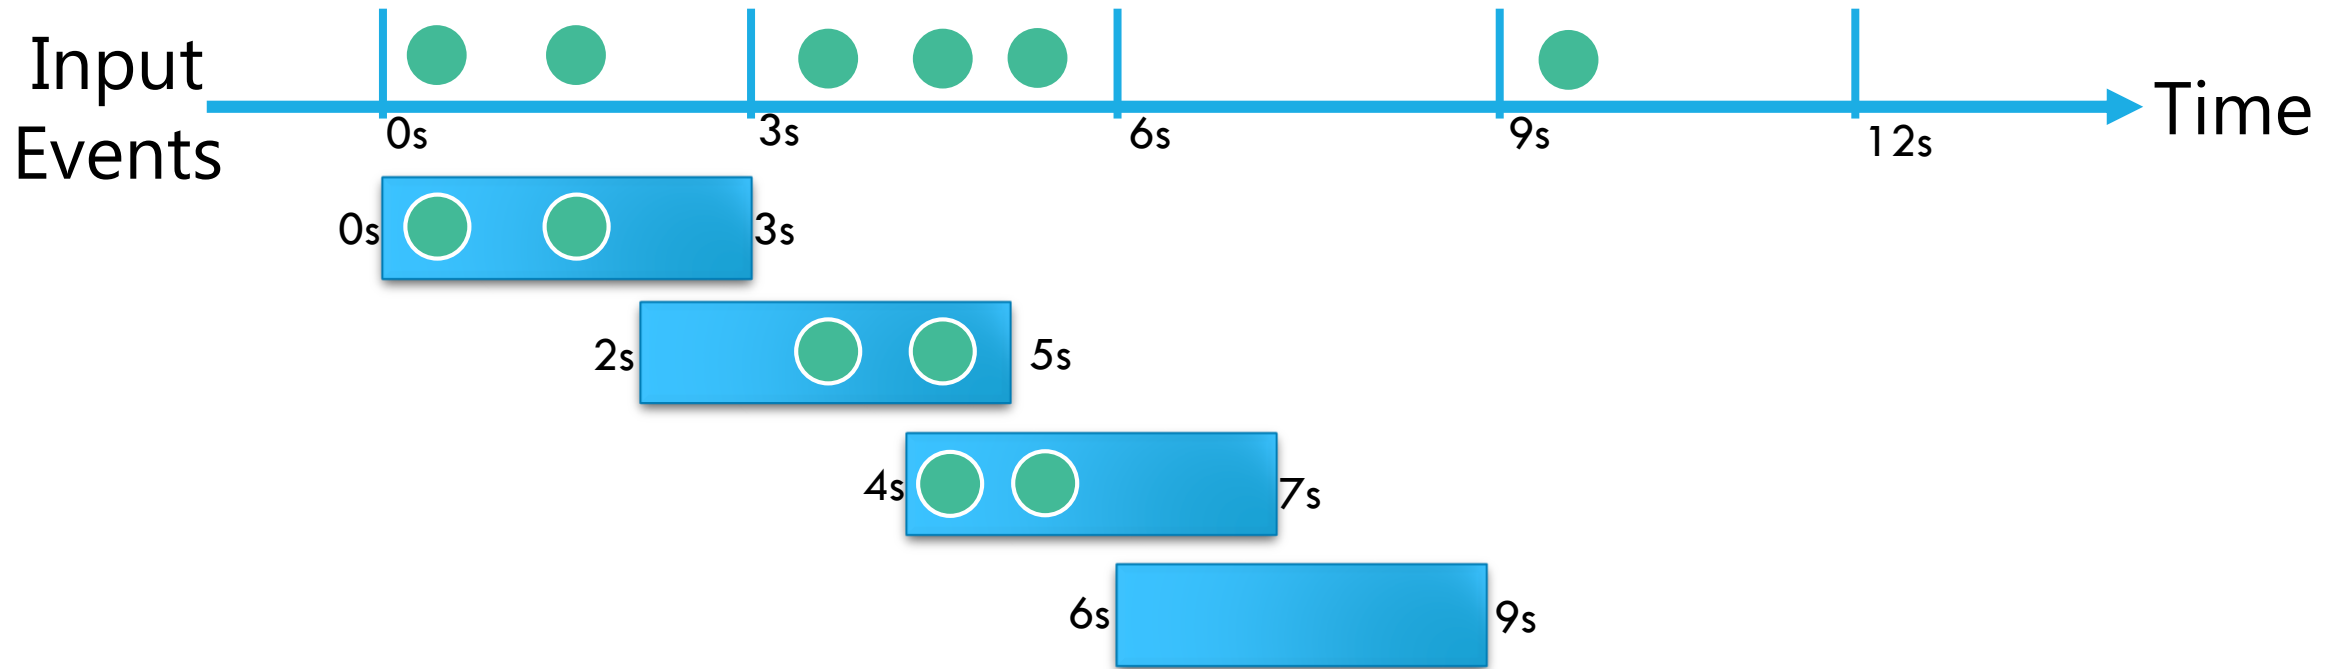
# TUMBLING WINDOW



```
SELECT TollID, System.Timestamp AS WindowEnd, Count(*) AS Count
FROM EntryStream TIMESTAMP BY EntryTime
GROUP BY TUMBLINGWINDOW(second, 3), TollID
```

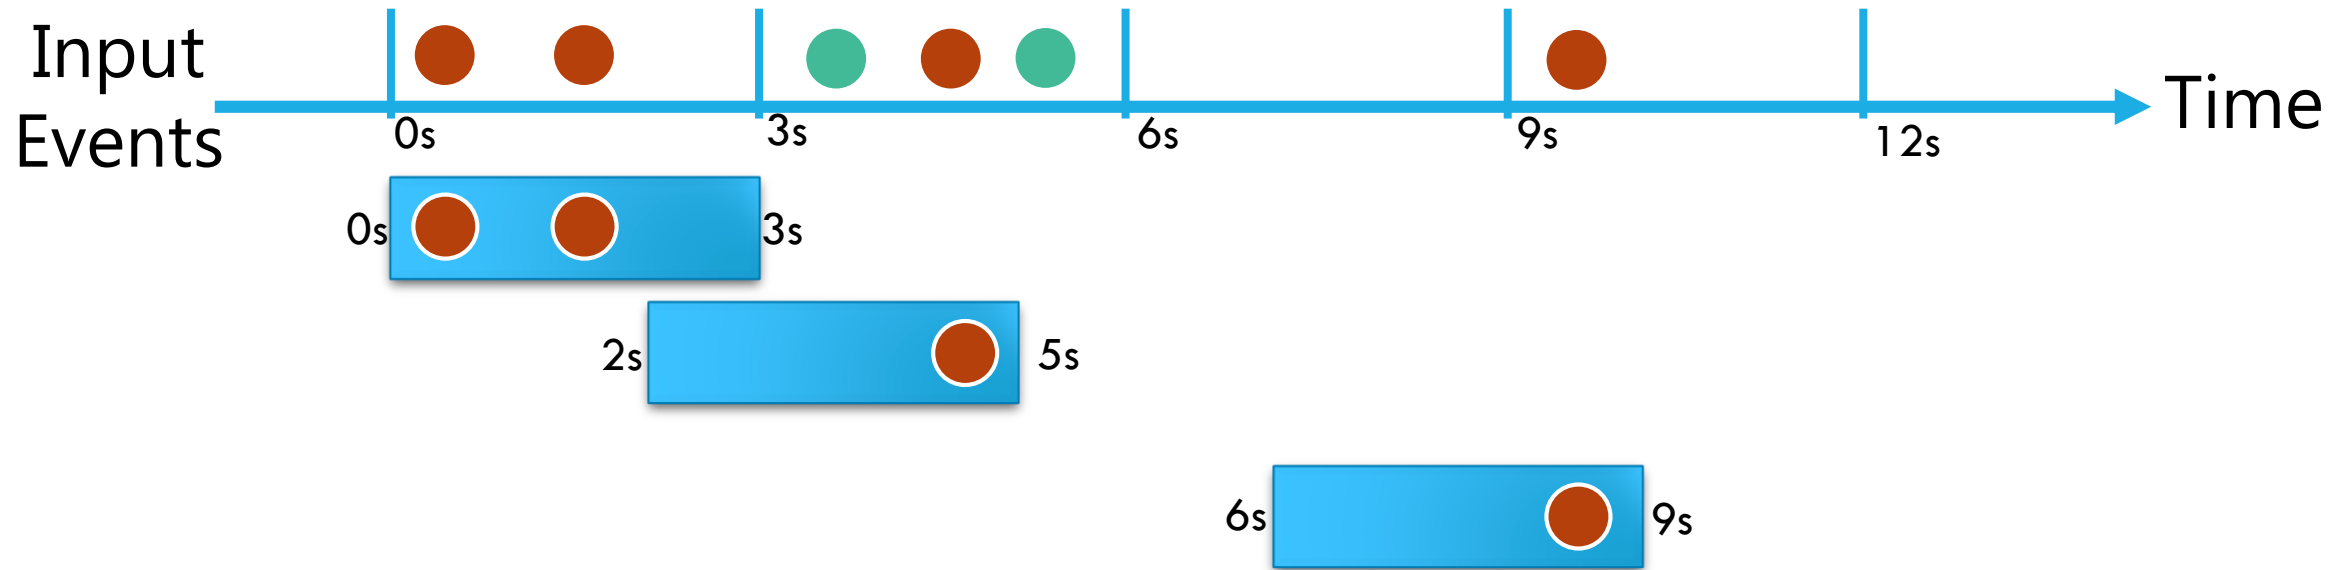


# HOPPING WINDOW



```
SELECT TollID, System.Timestamp AS WindowEnd, Count(*) AS Count
FROM EntryStream TIMESTAMP BY EntryTime
GROUP BY HOPPINGWINDOW(second, 2, 3), TollID
```

# SLIDING WINDOW



```
SELECT TollID, System.Timestamp AS WindowEnd, Count(*) AS Count
FROM EntryStream TIMESTAMP BY EntryTime
GROUP BY SLIDING(second, 2, 3), TollID
HAVING Color = RED
```

# SUM AGGREGATION

How much toll revenue is being accumulated every 3 minutes?

```
SELECT System.Timestamp AS WindowEnd, Sum(TollAmount) AS IntervalRevenue
FROM EntryStream TIMESTAMP BY EntryTime
GROUP BY TUMBLINGWINDOW(minute, 3), WindowEnd
```

# SUM AGGREGATION

Which 3 minute time interval made more than \$10?

```
SELECT System.Timestamp AS WindowEnd, Sum(TollAmount) AS IntervalRevenue
FROM EntryStream TIMESTAMP BY EntryTime
GROUP BY TUMBLINGWINDOW(minute, 3), WindowEnd
Having IntervalRevenue > 10
```



# DATEDIFF

How long does it take for each car to pass through the toll zone?

**SELECT**

EntryStream.LicensePlate,

**DATEDIFF** (**second**, EntryStream.EntryTime, Exitstream.ExitTime) **AS** DurationInSeconds

**FROM** EntryStream **timestamp** by EntryTime

**JOIN** Exitstream **timestamp** by ExitTime

**ON** Exitstream.LicensePlate = ExitStream.LicensePlate

**AND** **DATEDIFF**(**second**, EntryStream, ExitStream) **BETWEEN** 0 **AND** 1800

# DATEDIFF

How long does it take for each car to pass through the toll zone?

```
SELECT
    EntryStream.LicensePlate,
    DATEDIFF(second, EntryStream.EntryTime, Exitstream.ExitTime) AS DurationInSeconds
FROM EntryStream timestamp by EntryTime
JOIN Exitstream timestamp by ExitTime
ON Exitstream.LicensePlate = EntryStream.LicensePlate
AND DATEDIFF(second, EntryStream, ExitStream) BETWEEN 0 AND 1800
```

# DATEDIFF, INTEGER ONLY

How long (in HOURS) does it take for each car to pass through the toll zone?

- Decimal floats cut off, returns only 0s.

**SELECT**

**EntryStream.LicensePlate,**

**DATEDIFF(hour, EntryStream.EntryTime, Exitstream.ExitTime) AS DurationInSeconds**

**FROM EntryStream timestamp by EntryTime**

**JOIN Exitstream timestamp by ExitTime**

**ON Exitstream.LicensePlate = ExitStream.LicensePlate**

**AND DATEDIFF(hour, EntryStream, ExitStream) BETWEEN 0 AND 1800**

# CALCULATIONS

How fast (mph) was each car traveling through the toll zone?  
Assuming the toll zone was 1.5 miles long.

**SELECT**

EntryStream.LicensePlate,

1.5 / (**DATEDIFF**(**second**, (second, EntryStream.EntryTime, Exitstream.ExitTime) / 60 / 60) **AS MilesPerHour**

**FROM** EntryStream **timestamp** by EntryTime

**JOIN** Exitstream **timestamp** by ExitTime

**ON** Exitstream.LicensePlate = ExitStream.LicensePlate

**AND** **DATEDIFF**(**second**, EntryStream, ExitStream) **BETWEEN** 0 **AND** 3600

# AZURE STREAM-QL QUIRKS

Who was speeding through the toll zone?

- Simple question... but the query below will break.

**SELECT**

EntryStream.LicensePlate,

1.5 / (**DATEDIFF**(**second**, (second, EntryStream.EntryTime, Exitstream.ExitTime) / 60 / 60) **AS** MilesPerHour

**FROM** EntryStream **timestamp** by EntryTime

**JOIN** Exitstream **timestamp** by ExitTime

**ON** Exitstream.LicensePlate = ExitStream.LicensePlate

**AND** **DATEDIFF**(**second**, EntryStream, ExitStream) **BETWEEN** 0 **AND** 3600

~~**WHERE** MilesPerHour > 62~~



# AZURE STREAM-QL QUIRKS

Who was speeding through the toll zone?

- Works... but ugly solution.

**SELECT**

**EntryStream.LicensePlate,**

**1.5 / (DATEDIFF(second, (second, EntryStream.EntryTime, Exitstream.ExitTime) / 60 / 60) AS MilesPerHour**

**FROM EntryStream timestamp by EntryTime**

**JOIN Exitstream timestamp by ExitTime**

**ON Exitstream.LicensePlate = ExitStream.LicensePlate**

**AND DATEDIFF(second, EntryStream, ExitStream) BETWEEN 0 AND 3600**

**WHERE 1.5 / (DATEDIFF(second, (second, EntryStream.EntryTime, Exitstream.ExitTime) / 60 / 60) > 62**